

# Finalized\_Notebook

April 14, 2021

## 1 Credit Default Prediction: Code Notebook with Annotations

*Code is in Python using Jupyter Notebook*

*Exported using LaTeX*

*(n) indicates code reference was used, go to code references at the bottom to find references with links*

**#Table of contents**

---

### 1.1 1. Exploratory Data Analysis

- first look at data
- comparing defaulting and non-defaulting population
- basic graphing of data
- testing validity of data
- checking for problems that need to be addressed before fitting models

### 1.2 2. Data Preprocessing - Logistic Regression

- variable selection
- reducing collinearity in explanatory variables
- Normalization of data

### 1.3 3. Model Building and Testing - Logistic Regression

- Hyperparameter tuning
- Building and testing model
- Interpretation of model

### 1.4 4. Data Preprocessing - Random Forests

- variable selection
- missing value imputation

### 1.5 5. Model Building and Testing - Random Forests

- Hyperparameter tuning
- Building and testing model
- Interpretation of model

## 1.6 6. Comparing Models

- ROC curve comparison

## 1.7 7. Code References

- giving credit to people who helped us develop our code

# 2 Exploratory Data Analysis

Importing packages used in our analysis:

Numpy for scientific computing

Pandas for dealing with data tables

Matplotlib and seaborn for visualization

Statsmodels for calculating multicollinearity

Sklearn for building models and model interpretation, hyperparameter tuning

Autoimpute for imputation

```
[230]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sn
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.model_selection import RepeatedStratifiedKFold, GridSearchCV, \
    learning_curve, ShuffleSplit
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_curve, roc_auc_score
from autoimpute.imputations import SingleImputer
```

Importing all of our data sets:

```
[235]: df_test = pd.read_csv('Simulated_Data_Test.csv')
df_train = pd.read_csv('Simulated_Data_Train.csv')
df_validation = pd.read_csv('Simulated_Data_Validation.csv')
```

The first thing is to just get a basic summary of the train data we were given by looking at the raw data.

As we can see below, all variables except for state are of type float, and upon first glance it seems that the train data has imported correctly.

```
[240]: print(df_train.sample)
print(df_train.dtypes)
```

<bound method NDFrame.sample of			tot_credit_debt	avg_card_debt	
credit_age	credit_good_age	card_age	\		
0	80826.71	15872.99	300.0	114.0	292.0
1	96052.60	12178.02	281.0	102.0	232.0

2	75212.76	12052.24	261.0	149.0	260.0
3	70727.84	8416.80	227.0	93.0	223.0
4	41604.47	10611.97	249.0	136.0	241.0
...	...	...	...	...	...
19995	104765.01	13905.40	182.0	84.0	165.0
19996	83990.07	10325.02	320.0	129.0	280.0
19997	107606.69	17838.79	290.0	168.0	271.0
19998	78787.72	11447.61	208.0	104.0	194.0
19999	78296.90	10053.16	348.0	175.0	299.0

	non_mtg_acc_past_due_12_months_num	non_mtg_acc_past_due_6_months_num	\
0	0.0	0.0	
1	0.0	0.0	
2	0.0	0.0	
3	2.0	1.0	
4	0.0	0.0	
...	...	...	
19995	0.0	0.0	
19996	0.0	0.0	
19997	1.0	0.0	
19998	0.0	0.0	
19999	0.0	0.0	

	mortgages_past_due_6_months_num	credit_past_due_amount	\
0	0.0	0.00	
1	0.0	0.00	
2	0.0	0.00	
3	1.0	11013.96	
4	0.0	0.00	
...	...	...	
19995	0.0	0.00	
19996	0.0	0.00	
19997	0.0	0.00	
19998	0.0	0.00	
19999	0.0	0.00	

	inq_12_month_num	...	card_open_36_month_num	auto_open_36_month_num	\
0	3.0	...	0.0	0.0	
1	2.0	...	1.0	0.0	
2	1.0	...	0.0	1.0	
3	0.0	...	1.0	0.0	
4	0.0	...	0.0	0.0	
...	...	...	...	...	
19995	4.0	...	0.0	0.0	
19996	0.0	...	0.0	0.0	
19997	3.0	...	0.0	1.0	
19998	1.0	...	0.0	0.0	
19999	2.0	...	0.0	0.0	

	uti_card	uti_50plus_pct	uti_max_credit_line	uti_card_50plus_pct	\
0	0.365902	0.475594	0.410504		NaN
1	0.542786	0.543158	0.535147		0.587351
2	0.323678	0.321776	0.348713		0.413293
3	0.448721	0.422809	0.491365		0.466810
4	0.644030	0.619987	0.546655		0.588442
...	...	...	...	...	
19995	0.437699	0.557576	0.472592		0.481113
19996	0.637746	0.484714	0.734825		NaN
19997	0.486259	0.563475	0.406215		0.463120
19998	0.299035	0.283815	0.255758		0.281647
19999	0.574453	0.540931	0.540706		0.512513

	ind_acc_XYZ	rep_income	States	Default_ind
0	0.0	69000.0	AL	0.0
1	0.0	61000.0	FL	0.0
2	0.0	NaN	AL	0.0
3	0.0	79000.0	SC	1.0
4	1.0	NaN	LA	0.0
...	...	...	...	...
19995	1.0	NaN	GA	0.0
19996	0.0	87000.0	NC	0.0
19997	0.0	63000.0	SC	0.0
19998	0.0	71000.0	AL	0.0
19999	0.0	85000.0	AL	0.0

[20000 rows x 21 columns]>

tot_credit_debt	float64
avg_card_debt	float64
credit_age	float64
credit_good_age	float64
card_age	float64
non_mtg_acc_past_due_12_months_num	float64
non_mtg_acc_past_due_6_months_num	float64
mortgages_past_due_6_months_num	float64
credit_past_due_amount	float64
inq_12_month_num	float64
card_inq_24_month_num	float64
card_open_36_month_num	float64
auto_open_36_month_num	float64
uti_card	float64
uti_50plus_pct	float64
uti_max_credit_line	float64
uti_card_50plus_pct	float64
ind_acc_XYZ	float64
rep_income	float64
States	object

```
Default_ind          float64
dtype: object
```

Before we get ahead of ourselves, we want to check that there is nothing wrong with the test and validation datasets when compared to the train dataset.

Everything seems to be in order: correct data types, same number of variables, nothing seems out of the ordinary.

```
[236]: print(df_validation.sample)
       print(df_validation.dtypes)
```

```
<bound method NDFrame.sample of
credit_good_age  card_age  \
0          63651.27      9019.99      484.0      242.0      395.0
1          105559.29     16692.19      212.0      118.0      211.0
2           96062.99     10509.13      255.0      123.0      180.0
3           84417.40     13873.96      330.0      175.0      328.0
4          100623.91     15592.09      207.0      101.0      128.0
...
2995          90748.88     11481.81      274.0      161.0      268.0
2996         101930.98     15242.30      243.0      161.0      213.0
2997          74738.73     12175.49      351.0      179.0      271.0
2998         120357.58     14477.70      308.0      178.0      298.0
2999          74240.17     11873.45      341.0      181.0      333.0

non_mtg_acc_past_due_12_months_num  non_mtg_acc_past_due_6_months_num  \
0                                0.0                                0.0
1                                0.0                                0.0
2                                0.0                                0.0
3                                0.0                                0.0
4                                0.0                                0.0
...
2995                                0.0                                0.0
2996                                0.0                                0.0
2997                                0.0                                0.0
2998                                0.0                                0.0
2999                                0.0                                0.0

mortgages_past_due_6_months_num  credit_past_due_amount  \
0                                0.0                                0.0
1                                0.0                                0.0
2                                0.0                                0.0
3                                0.0                                0.0
4                                0.0                                0.0
...
2995                                0.0                                0.0
2996                                0.0                                0.0
2997                                0.0                                0.0
```

2998	0.0	0.0
2999	0.0	0.0

	inq_12_month_num	...	card_open_36_month_num	auto_open_36_month_num	\
0	0.0	...	0.0	0.0	
1	4.0	...	0.0	1.0	
2	2.0	...	0.0	0.0	
3	1.0	...	0.0	0.0	
4	6.0	...	0.0	0.0	
...	...	...	...	...	
2995	0.0	...	0.0	1.0	
2996	2.0	...	0.0	1.0	
2997	2.0	...	0.0	0.0	
2998	2.0	...	0.0	1.0	
2999	4.0	...	1.0	1.0	

	uti_card	uti_50plus_pct	uti_max_credit_line	uti_card_50plus_pct	\
0	0.619761	0.624652	0.506910	0.530109	
1	0.428082	0.533489	0.340476	0.388792	
2	0.437217	0.422643	0.418459	0.450523	
3	0.651360	0.566563	0.399319	0.562153	
4	0.586265	0.504849	0.652576	0.576409	
...	...	...	...	...	
2995	0.573039	0.665278	0.638087	0.559056	
2996	0.582873	0.662823	0.408222	0.564803	
2997	0.735002	0.586252	0.527243	0.651801	
2998	0.478272	0.545468	0.531942	0.557991	
2999	0.702814	0.658051	0.680245	0.736850	

	ind_acc_XYZ	rep_income	States	Default_ind
0	0.0	66000.0	FL	1.0
1	0.0	55000.0	MS	0.0
2	0.0	86000.0	MS	0.0
3	0.0	110000.0	MS	0.0
4	1.0	NaN	NC	0.0
...	...	...	...	...
2995	0.0	76000.0	FL	0.0
2996	0.0	65000.0	AL	0.0
2997	0.0	85000.0	LA	0.0
2998	1.0	NaN	NC	0.0
2999	0.0	47000.0	MS	0.0

[3000 rows x 21 columns]>

tot_credit_debt	float64
avg_card_debt	float64
credit_age	float64
credit_good_age	float64
card_age	float64

```

non_mtg_acc_past_due_12_months_num    float64
non_mtg_acc_past_due_6_months_num     float64
mortgages_past_due_6_months_num       float64
credit_past_due_amount                float64
inq_12_month_num                      float64
card_inq_24_month_num                 float64
card_open_36_month_num                float64
auto_open_36_month_num                float64
uti_card                              float64
uti_50plus_pct                        float64
uti_max_credit_line                   float64
uti_card_50plus_pct                   float64
ind_acc_XYZ                           float64
rep_income                            float64
States                                object
Default_ind                           float64
dtype: object

```

```

[238]: print(df_test.sample)
       print(df_test.dtypes)

```

```

<bound method NDFrame.sample of
credit_good_age  card_age  \
0          40477.81      7766.64      322.0      181.0      235.0
1          106760.98     16606.98      323.0      144.0      273.0
2          121428.34     13910.13      361.0      189.0      350.0
3           96515.05     15436.78      344.0      203.0      343.0
4          123760.22     14213.25      370.0      177.0      347.0
...          ...          ...          ...          ...          ...
4995         90788.44     11646.36      203.0      80.0      180.0
4996         91052.81     16505.20      291.0     122.0      285.0
4997         71061.71     11512.18      348.0     191.0      303.0
4998         82162.45     10014.90      357.0     163.0      339.0
4999        116943.32     18042.56      192.0      99.0      189.0

```

```

non_mtg_acc_past_due_12_months_num  non_mtg_acc_past_due_6_months_num  \
0                                0.0                                0.0
1                                0.0                                0.0
2                                0.0                                0.0
3                                0.0                                0.0
4                                0.0                                0.0
...                                ...                                ...
4995                             0.0                             0.0
4996                             0.0                             0.0
4997                             0.0                             0.0
4998                             0.0                             0.0
4999                             0.0                             0.0

```

	mortgages_past_due_6_months_num	credit_past_due_amount \
0	0.0	0.0
1	0.0	0.0
2	0.0	0.0
3	0.0	0.0
4	0.0	0.0
...	...	...
4995	0.0	0.0
4996	0.0	0.0
4997	0.0	0.0
4998	0.0	0.0
4999	0.0	0.0

	inq_12_month_num	...	card_open_36_month_num	auto_open_36_month_num \
0	2.0	...	0.0	1.0
1	2.0	...	0.0	0.0
2	0.0	...	0.0	0.0
3	0.0	...	0.0	0.0
4	0.0	...	0.0	0.0
...	...	...	...	...
4995	0.0	...	0.0	0.0
4996	3.0	...	0.0	0.0
4997	3.0	...	1.0	1.0
4998	1.0	...	0.0	0.0
4999	4.0	...	1.0	0.0

	uti_card	uti_50plus_pct	uti_max_credit_line	uti_card_50plus_pct \
0	0.588301	0.711887	0.538456	0.633699
1	0.707799	0.739732	0.562194	0.746901
2	0.479711	0.434104	0.374705	0.458438
3	0.443710	0.474768	0.372028	0.582352
4	0.571004	0.472634	0.605137	0.577697
...	...	...	...	...
4995	0.444254	0.629877	0.496559	0.398411
4996	0.330706	0.353832	0.485601	0.286623
4997	0.469755	0.478526	0.460937	0.443688
4998	0.089534	0.259846	0.127144	0.178131
4999	0.575263	0.593724	0.626940	0.503552

	ind_acc_XYZ	rep_income	States	Default_ind
0	0.0	26000.0	FL	0.0
1	1.0	47000.0	AL	0.0
2	0.0	71000.0	MS	0.0
3	1.0	68000.0	AL	0.0
4	0.0	76000.0	GA	0.0
...	...	...	...	...
4995	0.0	94000.0	LA	0.0
4996	1.0	73000.0	SC	0.0



4997	1.0	94000.0	LA	0.0
4998	0.0	70000.0	MS	0.0
4999	0.0	111000.0	MS	0.0

```
[5000 rows x 21 columns]>
tot_credit_debt          float64
avg_card_debt            float64
credit_age               float64
credit_good_age          float64
card_age                 float64
non_mtg_acc_past_due_12_months_num float64
non_mtg_acc_past_due_6_months_num float64
mortgages_past_due_6_months_num float64
credit_past_due_amount   float64
inq_12_month_num         float64
card_inq_24_month_num    float64
card_open_36_month_num   float64
auto_open_36_month_num   float64
uti_card                 float64
uti_50plus_pct           float64
uti_max_credit_line      float64
uti_card_50plus_pct      float64
ind_acc_XYZ              float64
rep_income               float64
States                   object
Default_ind              float64
dtype: object
```

Now we will garner summary stats about the all the variables. Important notes: - reported income mean is much higher than national mean. - 7.93% of population in train data set has defaulted - 25.855% of population already had an account with XYZ - All utilization variables have a mean around 50% - Standard deviation is relatively large for average card debt (0.66 as a proportion of the mean).

```
[49]: df_train.describe()
```

```
[49]:
```

	tot_credit_debt	avg_card_debt	credit_age	credit_good_age	\
count	20000.000000	20000.000000	20000.000000	20000.000000	
mean	94563.702530	14088.235475	296.697000	149.771750	
std	23546.443862	9314.495936	61.711702	34.016476	
min	2367.430000	2363.120000	54.000000	21.000000	
25%	78743.750000	11321.502500	255.000000	127.000000	
50%	94670.630000	13243.750000	297.000000	150.000000	
75%	110329.335000	15196.060000	339.000000	172.000000	
max	188890.960000	99999.000000	545.000000	296.000000	

	card_age	non_mtg_acc_past_due_12_months_num	\
count	20000.000000	20000.000000	

mean	268.015200	0.11135
std	59.364769	0.43389
min	41.000000	0.00000
25%	227.000000	0.00000
50%	268.000000	0.00000
75%	308.000000	0.00000
max	520.000000	4.00000

	non_mtg_acc_past_due_6_months_num	mortgages_past_due_6_months_num	\
count	20000.000000	20000.000000	
mean	0.027400	0.030200	
std	0.171903	0.171142	
min	0.000000	0.000000	
25%	0.000000	0.000000	
50%	0.000000	0.000000	
75%	0.000000	0.000000	
max	2.000000	1.000000	

	credit_past_due_amount	inq_12_month_num	card_inq_24_month_num	\
count	20000.000000	20000.000000	20000.000000	
mean	329.287867	1.762700	3.409600	
std	2073.899357	1.740816	2.926697	
min	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	1.000000	
50%	0.000000	1.000000	3.000000	
75%	0.000000	3.000000	5.000000	
max	32662.980000	10.000000	18.000000	

	card_open_36_month_num	auto_open_36_month_num	uti_card	\
count	20000.000000	20000.000000	20000.000000	
mean	0.163050	0.141000	0.503157	
std	0.386099	0.349607	0.109354	
min	0.000000	0.000000	0.065120	
25%	0.000000	0.000000	0.429611	
50%	0.000000	0.000000	0.502800	
75%	0.000000	0.000000	0.577412	
max	2.000000	2.000000	0.969289	

	uti_50plus_pct	uti_max_credit_line	uti_card_50plus_pct	ind_acc_XYZ	\
count	20000.000000	20000.000000	17945.000000	20000.000000	
mean	0.511007	0.507629	0.489594	0.258550	
std	0.113456	0.108624	0.119701	0.437849	
min	0.033749	0.005174	0.000000	0.000000	
25%	0.435171	0.433550	0.409794	0.000000	
50%	0.509922	0.507193	0.490074	0.000000	
75%	0.588418	0.581376	0.569036	1.000000	
max	0.988964	1.000000	0.970776	1.000000	

	rep_income	Default_ind
count	18430.000000	20000.000000
mean	75499.511666	0.079300
std	16361.955146	0.270213
min	12000.000000	0.000000
25%	64000.000000	0.000000
50%	75000.000000	0.000000
75%	86000.000000	0.000000
max	150000.000000	1.000000

We want to check which values are missing.

At this point, we are considering dropping the reported income variable due to its self reported nature.

We will need to find a way to deal with uti\_card\_50plus\_pct as that may be an important variable.

```
[23]: print(df_train.isnull().sum())
      print(df_train[['uti_card_50plus_pct', 'rep_income']].isnull().sum())
```

```
tot_credit_debt          0
avg_card_debt            0
credit_age               0
credit_good_age          0
card_age                 0
non_mtg_acc_past_due_12_months_num  0
non_mtg_acc_past_due_6_months_num  0
mortgages_past_due_6_months_num    0
credit_past_due_amount    0
inq_12_month_num         0
card_inq_24_month_num     0
card_open_36_month_num    0
auto_open_36_month_num    0
uti_card                 0
uti_50plus_pct           0
uti_max_credit_line       0
uti_card_50plus_pct      2055
ind_acc_XYZ              0
rep_income               1570
AL                        0
FL                        0
GA                        0
LA                        0
MS                        0
NC                        0
SC                        0
```

```

Default_ind                                0
dtype: int64
uti_card_50plus_pct      2055
rep_income                1570
dtype: int64

```

It is important to compare the statistics of the defaulting population to the non defaulting population.

We note the following: - The mean average card debt, credit past due amount, and utilization variables are much higher in the defaulting population than the non-defaulting population - For average card debt and credit past due amount, the defaulting population had a much greater std, suggesting high variance in the defaulting population - Reported income stats for both populations is very similar

```

[8]: defaulting_population = df_train[df_train["Default_ind"] == 1]
      non_defaulting_population = df_train[df_train["Default_ind"] == 0]

```

```

[9]: defaulting_population.describe()

```

```

[9]:      tot_credit_debt  avg_card_debt  credit_age  credit_good_age  \
count      1586.000000      1586.000000  1586.000000      1586.000000
mean      91915.567554     17880.428821    275.160151      140.535309
std       29920.362673     21195.058690     61.986033       33.557997
min        6898.500000      2363.120000     54.000000       21.000000
25%       70589.517500      9864.232500    233.000000      118.000000
50%       91235.590000     13125.530000    274.000000      140.000000
75%      112973.610000     16356.372500    316.000000      163.000000
max      188890.960000     99999.000000    521.000000      274.000000

      card_age  non_mtg_acc_past_due_12_months_num  \
count      1586.000000      1586.000000
mean       247.708701         0.706810
std         58.602115         1.035236
min         41.000000         0.000000
25%        206.000000         0.000000
50%        247.000000         0.000000
75%        288.000000         2.000000
max        463.000000         4.000000

      non_mtg_acc_past_due_6_months_num  mortgages_past_due_6_months_num  \
count      1586.000000      1586.000000
mean         0.249685         0.268600
std         0.473343         0.443371
min          0.000000         0.000000
25%          0.000000         0.000000
50%          0.000000         0.000000
75%          0.000000         1.000000

```

max	2.000000	1.000000
-----	----------	----------

	credit_past_due_amount	inq_12_month_num	card_inq_24_month_num \
count	1586.000000	1586.000000	1586.000000
mean	3060.922629	2.288777	4.163934
std	5777.479779	1.867494	3.137177
min	0.000000	0.000000	0.000000
25%	0.000000	1.000000	2.000000
50%	0.000000	2.000000	4.000000
75%	4905.157500	4.000000	6.000000
max	32662.980000	9.000000	15.000000

	card_open_36_month_num	auto_open_36_month_num	uti_card \
count	1586.000000	1586.000000	1586.000000
mean	0.192938	0.149433	0.560360
std	0.419524	0.356627	0.110701
min	0.000000	0.000000	0.208749
25%	0.000000	0.000000	0.487653
50%	0.000000	0.000000	0.564220
75%	0.000000	0.000000	0.636084
max	2.000000	1.000000	0.969289

	uti_50plus_pct	uti_max_credit_line	uti_card_50plus_pct	ind_acc_XYZ \
count	1586.000000	1586.000000	1412.000000	1586.000000
mean	0.557966	0.551171	0.544364	0.206810
std	0.114920	0.109630	0.123059	0.405145
min	0.163960	0.186306	0.118153	0.000000
25%	0.476796	0.476078	0.462664	0.000000
50%	0.558708	0.550171	0.547647	0.000000
75%	0.638486	0.623271	0.629132	0.000000
max	0.894996	1.000000	0.970776	1.000000

	rep_income	Default_ind
count	1457.000000	1586.0
mean	74522.992450	1.0
std	16775.126303	0.0
min	26000.000000	1.0
25%	63000.000000	1.0
50%	74000.000000	1.0
75%	85000.000000	1.0
max	123000.000000	1.0

```
[242]: non_defaulting_population.describe()
```

```
[242]:
```

	tot_credit_debt	avg_card_debt	credit_age	credit_good_age \
count	18414.000000	18414.000000	18414.000000	18414.000000
mean	94791.786710	13761.613413	298.551971	150.567286

std	22901.385098	7363.255055	61.336991	33.939197
min	2367.430000	4595.020000	78.000000	27.000000
25%	79292.612500	11401.827500	257.000000	128.000000
50%	94912.085000	13251.240000	299.000000	151.000000
75%	110188.782500	15138.300000	340.000000	173.000000
max	182858.990000	99999.000000	545.000000	296.000000

	card_age	non_mtg_acc_past_due_12_months_num	\
count	18414.000000	18414.000000	
mean	269.764201	0.060063	
std	59.106129	0.281161	
min	56.000000	0.000000	
25%	229.000000	0.000000	
50%	270.000000	0.000000	
75%	310.000000	0.000000	
max	520.000000	3.000000	

	non_mtg_acc_past_due_6_months_num	mortgages_past_due_6_months_num	\
count	18414.000000	18414.000000	
mean	0.008255	0.009667	
std	0.090481	0.097845	
min	0.000000	0.000000	
25%	0.000000	0.000000	
50%	0.000000	0.000000	
75%	0.000000	0.000000	
max	1.000000	1.000000	

	credit_past_due_amount	inq_12_month_num	card_inq_24_month_num	\
count	18414.000000	18414.000000	18414.000000	
mean	94.011842	1.717389	3.344629	
std	1048.877699	1.722024	2.898779	
min	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	1.000000	
50%	0.000000	1.000000	3.000000	
75%	0.000000	3.000000	5.000000	
max	24211.550000	10.000000	18.000000	

	card_open_36_month_num	auto_open_36_month_num	uti_card	\
count	18414.000000	18414.000000	18414.000000	
mean	0.160476	0.140274	0.498230	
std	0.382987	0.348996	0.107830	
min	0.000000	0.000000	0.065120	
25%	0.000000	0.000000	0.425885	
50%	0.000000	0.000000	0.498045	
75%	0.000000	0.000000	0.571530	
max	2.000000	2.000000	0.922326	

	uti_50plus_pct	uti_max_credit_line	uti_card_50plus_pct	ind_acc_XYZ \
count	18414.000000	18414.000000	16533.000000	18414.000000
mean	0.506963	0.503879	0.484916	0.263006
std	0.112418	0.107720	0.118243	0.440278
min	0.033749	0.005174	0.000000	0.000000
25%	0.431793	0.430766	0.405870	0.000000
50%	0.506286	0.503482	0.485286	0.000000
75%	0.583785	0.577646	0.563275	1.000000
max	0.988964	0.971640	0.949959	1.000000

	rep_income	Default_ind
count	16973.000000	18414.0
mean	75583.338243	0.0
std	16323.783467	0.0
min	12000.000000	0.0
25%	65000.000000	0.0
50%	76000.000000	0.0
75%	87000.000000	0.0
max	150000.000000	0.0

Do we have enough observations for the amount of explanatory variables? One way to check this for regression analysis is the rule of ten test. If  $10k/p > n$ , there are too many explanatory variables and not enough observations. -  $k$  = explanatory variables = 20 -  $p$  = probability that account defaulted = 0.0793 (found in summary stats) -  $n$  = number of observations = 20,000

Data has passed rule of ten test since 20,000 is much larger than 2522.

For random forest, the number of observations is less of a concern for us as random forests are good with both large and small data sets.

```
[201]: print((10 * 20) / 0.0793)
```

```
2522.068095838588
```

We can check for multicollinearity by finding each variable's variance inflation factor. This is extremely important to consider for the logistic regression model, as one of the main assumptions is that there is little to no collinearity explanatory variables.

As can be seen, there are many VIF values which are greater than 10, an unacceptable result. Additionally, there some VIF values that are less than 10 but greater than 5, a cause for concern (1).

```
[17]: def calc_vif(X):
        vif = pd.DataFrame()
        vif["variables"] = X.columns
        vif["VIF"] = [variance_inflation_factor(X.values, i) for i in range(X.
↪shape[1])]

```

```

    return(vif)

columns = []
for num in range(0,18):
    if num == 16:
        pass
    else:
        columns.append(num)

calc_vif(df_train.iloc[:,columns])

```

```

[17]:
      variables      VIF
0      tot_credit_debt  16.224066
1      avg_card_debt   3.835799
2      credit_age     232.525965
3      credit_good_age  52.883615
4      card_age       174.509147
5  non_mtg_acc_past_due_12_months_num  4.612476
6  non_mtg_acc_past_due_6_months_num   3.462795
7  mortgages_past_due_6_months_num    6.670029
8      credit_past_due_amount   6.570714
9      inq_12_month_num    7.769188
10     card_inq_24_month_num    9.002325
11     card_open_36_month_num    1.188654
12     auto_open_36_month_num    1.172338
13           uti_card    78.503559
14     uti_50plus_pct    46.474318
15     uti_max_credit_line    49.135164
16     ind_acc_XYZ     1.348725

```

Let's check for collinearity between specific variables, including the response variable. This will help shed light on which explanatory variables are colinear, which will be extremely important for logistic regression. We will also discover which explanatory variables we may consider to have too low of a correlation to the default indicator, which we want to minimize for our random tree model.

To create a correlation matrix of all variables, we will firstly create dummy variables for the states.

```

[204]: dummies = pd.get_dummies(df_train['States'])
default_ind = df_train['Default_ind']
df_train = df_train.drop(columns=['States', 'Default_ind'])
df_train = df_train.join(dummies)
df_train = df_train.join(default_ind)

```

Now the data is ready for a correlation matrix.

Logistic regression notes:




It seems that reported income does not have a high correlation with the response variable. This confirms to us that it is no longer worth keeping. The first five columns present unacceptable levels of colinearity. With no easy way to combine them, the four with least amount of correlation to the response variable will be dropped. Since the 12 month and 24 month card inquiry variables are highly colinear and since the 24 month card inquiry variable is cumulative, we will drop the 12 month card inquiry variable.

Random forest notes:

For now, we will consider credit age, credit good age, and card age to have too low of a correlation to the default indicator and disregard them for our random tree model.

```
[22]: columns = []
fig, ax = plt.subplots(figsize=(12,10))
for num in range(0,27):
    columns.append(num)
corrMatrix = df_train.iloc[:,columns].corr()
sn.heatmap(corrMatrix, annot=False, ax=ax)
plt.show()
```




For logistic regression, we will need to deal with missing values, as the model will not do this for us. Since the only variable with missing values we decided to keep is `uti_card_50plus_pct`, we will need to find a way to deal with missing values for this column. Our first idea is, since utilization variables have high collinearity and are on the same scale, to combine all utilization variables into a mean utilization variable which will consider `uti_card_50plus_pct` only if available for the observation. To make sure this is feasible, we must first check the normality of the utilization variables.

As can be seen below, the distribution of all utilization variables is approximately normal, meaning that mean utilization could be a useful variable.


```
[13]: sn.displot(df_train,x="uti_card")
```

```
[13]: <seaborn.axisgrid.FacetGrid at 0x2126627ed30>
```




```
[18]: sn.displot(df_train,x="uti_card_50plus_pct")
```

```
[18]: <seaborn.axisgrid.FacetGrid at 0x212669e4a00>
```




```
[15]: sn.displot(df_train,x="uti_max_credit_line")
```

```
[15]: <seaborn.axisgrid.FacetGrid at 0x212669ee790>
```



```
[19]: sn.displot(df_train,x="uti_50plus_pct")
```

```
[19]: <seaborn.axisgrid.FacetGrid at 0x212672312e0>
```



### 3 Data Preprocessing - Logistic Regression

We have already determined based on EDA which explanatory variables need to be dropped: tot\_credit\_debt, credit\_age, credit\_good\_age, card\_age, rep\_income, inq\_12\_month\_num.

```
[205]: df_train = df_train.drop(columns=["tot_credit_debt", "credit_age",
    ↪ "credit_good_age", "card_age", "rep_income", "inq_12_month_num"])
```

Next comes the process of combining all utilization variables into one variable, and replacing all the instances with the new overlapping variable. By taking the average utilization, we are able to preserve the available data in uti\_card\_50plus\_pct.

```
[206]: util_vars = ["uti_card", "uti_50plus_pct", "uti_max_credit_line",
    ↪ "uti_card_50plus_pct"]
util_credit = df_train[util_vars]
util_credit["mean"] = util_credit.mean(axis=1)
```

```
df_train = df_train.drop(columns=["uti_50plus_pct", "uti_max_credit_line",
    ↪ "uti_card_50plus_pct"])
df_train["uti_card"] = util_credit["mean"]
df_train = df_train.rename(columns={"uti_card": "mean_uti"})
```

<ipython-input-206-543dfd05b1c6>:3: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
`util_credit["mean"] = util_credit.mean(axis=1)`

Past due variables have unacceptable levels of colinearity. It makes sense to combine the non mortgage and mortgages past due 6 month variables as they span the same time period and are highly colinear. For now, the other past due variables are not dropped due to their high correlation with the response variable.

```
[207]: past_due_vars = ["non_mtg_acc_past_due_6_months_num",
    ↪ "mortgages_past_due_6_months_num"]
past_due_credit = df_train[past_due_vars]
past_due_credit["past_due_6_months_num"] =
    ↪ past_due_credit["non_mtg_acc_past_due_6_months_num"] +
    ↪ past_due_credit["mortgages_past_due_6_months_num"]

df_train = df_train.drop(columns=["mortgages_past_due_6_months_num"])
df_train["non_mtg_acc_past_due_6_months_num"] =
    ↪ past_due_credit["past_due_6_months_num"]
df_train = df_train.rename(columns={"non_mtg_acc_past_due_6_months_num":
    ↪ "past_due_6_months_num"})
```

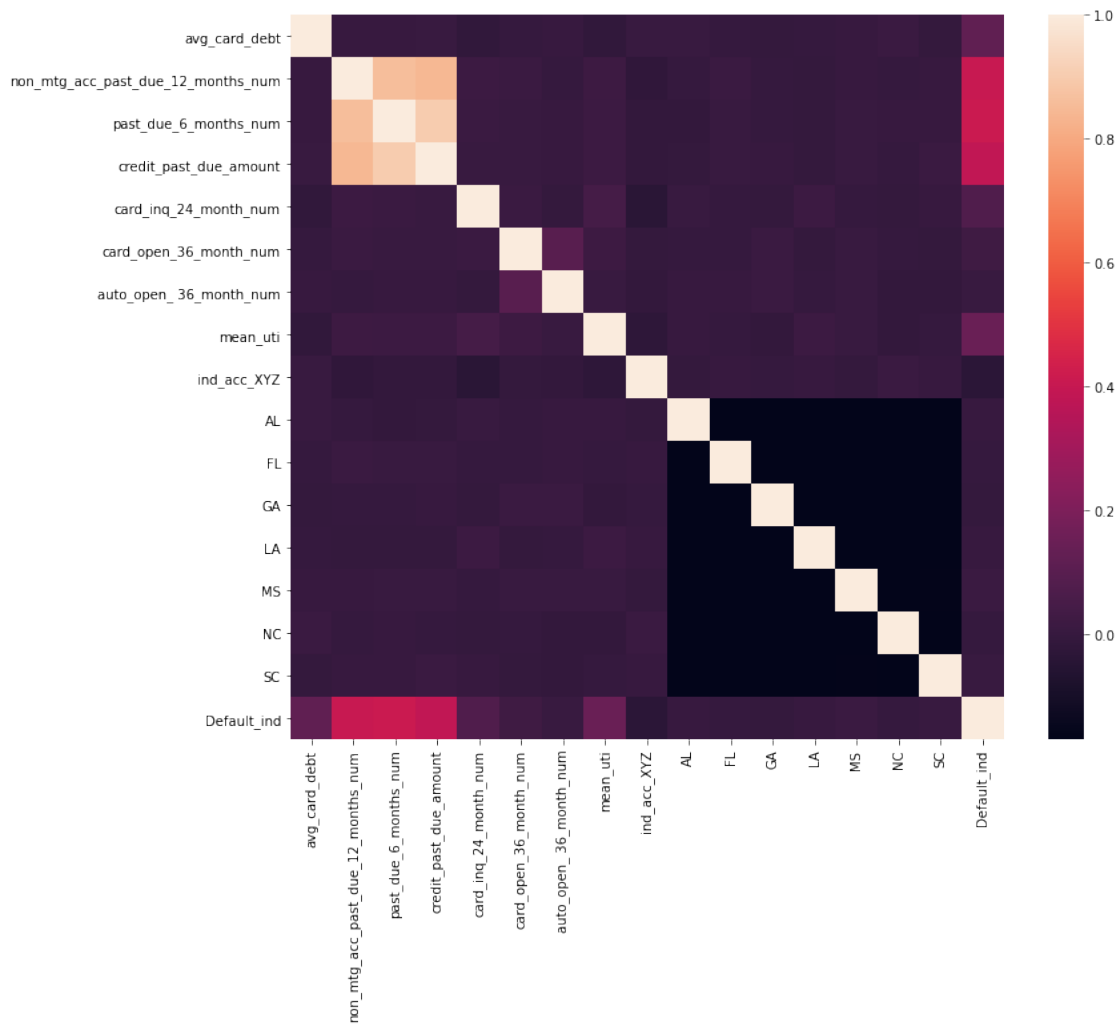
<ipython-input-207-db6338c54e88>:3: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
`past_due_credit["past_due_6_months_num"] =  
past_due_credit["non_mtg_acc_past_due_6_months_num"] +  
past_due_credit["mortgages_past_due_6_months_num"]`

We will check collinearity and multicollinearity again to see if we have reduced it to an acceptable level.

While colinearity between the past due variables' colinearity is worrying, they are simply too valuable to drop for now due to their high linearity with the response variable. The VIF values give an encouraging picture, however, as values, while slightly high, are in an acceptable range (1).

```
[31]: columns = []
fig, ax = plt.subplots(figsize=(12,10))
for num in range(0,17):
    columns.append(num)
corrMatrix = df_train.iloc[:,columns].corr()
sn.heatmap(corrMatrix, annot=False, ax=ax)
plt.show()
```



```
[33]: columns = []
for num in range(0,16):
    columns.append(num)

calc_vif(df_train.iloc[:,columns])
```

```
[33]:
```

	variables	VIF
0	avg_card_debt	1.000904
1	non_mtg_acc_past_due_12_months_num	4.245805
2	past_due_6_months_num	6.522983
3	credit_past_due_amount	5.945257
4	card_inq_24_month_num	1.004331
5	card_open_36_month_num	1.010175
6	auto_open_36_month_num	1.009894
7	mean_util	1.004178
8	ind_acc_XYZ	1.002969
9	AL	5.388244
10	FL	5.303679
11	GA	5.263706
12	LA	5.365403
13	MS	5.287676
14	NC	5.340952
15	SC	5.242121

Last thing left to do is normalize data:

```
[208]: df_train_norm=((df_train-df_train.min())/(df_train.max()-df_train.min()))*20
```

Check preprocessed data:

Everything seems to be in order. Time to move on to model building.

```
[36]: print(df_train_norm.sample)
print(df_train_norm.shape)
print(df_train_norm.dtypes)
```

```
<bound method NDFrame.sample of
non_mtg_acc_past_due_12_months_num \
0          2.767399          0.0
1          2.010511          0.0
2          1.984746          0.0
3          1.240052         10.0
4          1.689717          0.0
...          ...          ...
19995        2.364352          0.0
19996        1.630937          0.0
19997        3.170078          5.0
19998        1.860892          0.0
19999        1.575249          0.0

      past_due_6_months_num  credit_past_due_amount  card_inq_24_month_num \
0          0.000000          0.000000          4.444444
1          0.000000          0.000000          4.444444
2          0.000000          0.000000          3.333333
3         13.333333          6.744002          1.111111
```



4	0.000000	0.000000	2.222222
...	...	...	...
19995	0.000000	0.000000	4.444444
19996	0.000000	0.000000	1.111111
19997	0.000000	0.000000	3.333333
19998	0.000000	0.000000	2.222222
19999	0.000000	0.000000	4.444444

	card_open_36_month_num	auto_open_36_month_num	mean_util	\
0	0.0	0.0	7.470849	
1	10.0	0.0	10.830203	
2	0.0	10.0	5.839036	
3	10.0	0.0	8.470172	
4	0.0	0.0	12.018333	
...	...	...	...	
19995	0.0	0.0	9.213416	
19996	0.0	0.0	12.499806	
19997	0.0	10.0	9.027027	
19998	0.0	0.0	4.049374	
19999	0.0	0.0	10.581955	

	ind_acc_XYZ	AL	FL	GA	LA	MS	NC	SC	Default_ind
0	0.0	20.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
1	0.0	0.0	20.0	0.0	0.0	0.0	0.0	0.0	0.0
2	0.0	20.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	20.0	20.0
4	20.0	0.0	0.0	0.0	20.0	0.0	0.0	0.0	0.0
...	...	...	...	...	...	...	...	...	...
19995	20.0	0.0	0.0	20.0	0.0	0.0	0.0	0.0	0.0
19996	0.0	0.0	0.0	0.0	0.0	0.0	20.0	0.0	0.0
19997	0.0	0.0	0.0	0.0	0.0	0.0	0.0	20.0	0.0
19998	0.0	20.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
19999	0.0	20.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

[20000 rows x 17 columns]>

(20000, 17)

avg_card_debt	float64
non_mtg_acc_past_due_12_months_num	float64
past_due_6_months_num	float64
credit_past_due_amount	float64
card_inq_24_month_num	float64
card_open_36_month_num	float64
auto_open_36_month_num	float64
mean_util	float64
ind_acc_XYZ	float64
AL	float64
FL	float64
GA	float64

LA	float64
MS	float64
NC	float64
SC	float64
Default_ind	float64
dtype:	object

## 4 Model Building and Testing - Logistic Regression

We will begin by performing hyper parameter optimization on the validation data set for logistic regression. In order to do this, we must manipulate the validation data set in the same way we have manipulated the training data set.

```
[209]: dummies = pd.get_dummies(df_validation['States'])
default_ind = df_validation['Default_ind']
df_validation = df_validation.
    ↳drop(columns=["Default_ind", "States", "tot_credit_debt", "credit_age",
    ↳"credit_good_age", "card_age", "rep_income", "inq_12_month_num"])
df_validation = df_validation.join(dummies)
df_validation = df_validation.join(default_ind)
util_vars = ["uti_card", "uti_50plus_pct", "uti_max_credit_line",
    ↳"uti_card_50plus_pct"]
util_credit = df_validation[util_vars]
util_credit["mean"] = util_credit.mean(axis=1)
df_validation = df_validation.drop(columns=["uti_50plus_pct",
    ↳"uti_max_credit_line", "uti_card_50plus_pct"])
df_validation["uti_card"] = util_credit["mean"]
df_validation = df_validation.rename(columns={"uti_card": "mean_uti"})
past_due_vars = ["non_mtg_acc_past_due_6_months_num",
    ↳"mortgages_past_due_6_months_num"]
past_due_credit = df_validation[past_due_vars]
past_due_credit["past_due_6_months_num"] =
    ↳past_due_credit["non_mtg_acc_past_due_6_months_num"] +
    ↳past_due_credit["mortgages_past_due_6_months_num"]
df_validation = df_validation.drop(columns=["mortgages_past_due_6_months_num"])
df_validation["non_mtg_acc_past_due_6_months_num"] =
    ↳past_due_credit["past_due_6_months_num"]
df_validation = df_validation.
    ↳rename(columns={"non_mtg_acc_past_due_6_months_num": "past_due_6_months_num"})
df_validation_norm = ((df_validation - df_validation.min()) / (df_validation.
    ↳max() - df_validation.min())) * 20
```

```
<ipython-input-209-959fbd1f3629>:8: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <https://pandas.pydata.org/pandas->

```
docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
    util_credit["mean"] = util_credit.mean(axis=1)
<ipython-input-209-959fbd1f3629>:14: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
    past_due_credit["past_due_6_months_num"] =
past_due_credit["non_mtg_acc_past_due_6_months_num"] +
past_due_credit["mortgages_past_due_6_months_num"]
```

Everything looks in order:

```
[39]: print(df_validation_norm.sample)
      print(df_validation_norm.shape)
      print(df_validation_norm.dtypes)
```

```
<bound method NDFrame.sample of          avg_card_debt
non_mtg_acc_past_due_12_months_num  \
0          1.258527          0.0
1          2.838983          0.0
2          1.565286          0.0
3          2.258434          0.0
4          2.612365          0.0
...          ...          ...
2995         1.765656          0.0
2996         2.540309          0.0
2997         1.908553          0.0
2998         2.382803          0.0
2999         1.846333          0.0
```

```
      past_due_6_months_num  credit_past_due_amount  card_inq_24_month_num  \
0          0.0          0.0          0.000000
1          0.0          0.0         10.769231
2          0.0          0.0          4.615385
3          0.0          0.0          1.538462
4          0.0          0.0         13.846154
...          ...          ...          ...
2995         0.0          0.0          0.000000
2996         0.0          0.0          0.000000
2997         0.0          0.0          6.153846
2998         0.0          0.0          3.076923
2999         0.0          0.0         10.769231
```

```
      card_open_36_month_num  auto_open_ 36_month_num  mean_util  ind_acc_XYZ  \
0          0.0          0.0      14.171756          0.0
1          0.0         10.0       9.861754          0.0
2          0.0          0.0      10.139092          0.0
```

3	0.0	0.0	13.427118	0.0
4	0.0	0.0	14.453941	20.0
...	...	...	...	...
2995	0.0	10.0	15.295818	0.0
2996	0.0	10.0	13.714113	0.0
2997	0.0	0.0	15.768993	0.0
2998	0.0	10.0	12.947497	20.0
2999	10.0	10.0	17.795304	0.0

	AL	FL	GA	LA	MS	NC	SC	Default_ind
0	0.0	20.0	0.0	0.0	0.0	0.0	0.0	20.0
1	0.0	0.0	0.0	0.0	20.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	20.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	20.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0	20.0	0.0	0.0
...	...	...	...	...	...	...	...	...
2995	0.0	20.0	0.0	0.0	0.0	0.0	0.0	0.0
2996	20.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2997	0.0	0.0	0.0	20.0	0.0	0.0	0.0	0.0
2998	0.0	0.0	0.0	0.0	0.0	20.0	0.0	0.0
2999	0.0	0.0	0.0	0.0	20.0	0.0	0.0	0.0

```
[3000 rows x 17 columns]>
```

```
(3000, 17)
```

```
avg_card_debt                float64
non_mtg_acc_past_due_12_months_num  float64
past_due_6_months_num          float64
credit_past_due_amount         float64
card_inq_24_month_num          float64
card_open_36_month_num         float64
auto_open_36_month_num         float64
mean_util                      float64
ind_acc_XYZ                    float64
AL                             float64
FL                             float64
GA                             float64
LA                             float64
MS                             float64
NC                             float64
SC                             float64
Default_ind                    float64
dtype: object
```

We will now perform hyperparameter tuning using our validation data set (2a).

The parameters we have chosen are: - solvers: newtown conjugate gradient, limited memory bfgs, and liblinear - penalty: l2 only (since it works for all three solvers) - C values: 100, 10, 1.0, 0.1, 0.01

```
[210]: model = LogisticRegression()
solvers = ['newton-cg', 'lbfgs', 'liblinear']
penalty = ['l2']
c_values = [100, 10, 1.0, 0.1, 0.01]
grid = dict(solver=solvers,penalty=penalty,C=c_values)
cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
grid_search = GridSearchCV(estimator=model, param_grid=grid, n_jobs=-1, cv=cv,
    ↳scoring='accuracy',error_score=0)
grid_result = grid_search.fit(df_validation_norm.iloc[:,range(0,16)],
    ↳df_validation_norm.iloc[:,16])
print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))
```

```
Best: 0.937778 using {'C': 0.1, 'penalty': 'l2', 'solver': 'liblinear'}
0.937667 (0.009195) with: {'C': 100, 'penalty': 'l2', 'solver': 'newton-cg'}
0.937667 (0.009195) with: {'C': 100, 'penalty': 'l2', 'solver': 'lbfgs'}
0.937667 (0.009195) with: {'C': 100, 'penalty': 'l2', 'solver': 'liblinear'}
0.937667 (0.009195) with: {'C': 10, 'penalty': 'l2', 'solver': 'newton-cg'}
0.937667 (0.009195) with: {'C': 10, 'penalty': 'l2', 'solver': 'lbfgs'}
0.937667 (0.009195) with: {'C': 10, 'penalty': 'l2', 'solver': 'liblinear'}
0.937667 (0.009195) with: {'C': 1.0, 'penalty': 'l2', 'solver': 'newton-cg'}
0.937667 (0.009195) with: {'C': 1.0, 'penalty': 'l2', 'solver': 'lbfgs'}
0.937667 (0.009195) with: {'C': 1.0, 'penalty': 'l2', 'solver': 'liblinear'}
0.937667 (0.009195) with: {'C': 0.1, 'penalty': 'l2', 'solver': 'newton-cg'}
0.937667 (0.009195) with: {'C': 0.1, 'penalty': 'l2', 'solver': 'lbfgs'}
0.937778 (0.009081) with: {'C': 0.1, 'penalty': 'l2', 'solver': 'liblinear'}
0.937778 (0.008664) with: {'C': 0.01, 'penalty': 'l2', 'solver': 'newton-cg'}
0.937778 (0.008664) with: {'C': 0.01, 'penalty': 'l2', 'solver': 'lbfgs'}
0.936667 (0.008300) with: {'C': 0.01, 'penalty': 'l2', 'solver': 'liblinear'}
```

Based on this, the most accurate model has the following hyperparameters: - solver: liblinear - penalty: l2 - C value: 0.1

Train model based on chosen hyperparameters using the train data set:

```
[211]: lrModel = LogisticRegression(C = 0.1, penalty = "l2",solver="liblinear")
lrModel = lrModel.fit(df_train_norm.iloc[:,range(0,16)], df_train_norm.iloc[:
    ↳,16])
```

Transforming test data to match training data:

```
[212]: dummies = pd.get_dummies(df_test['States'])
default_ind = df_test['Default_ind']
df_test = df_test.drop(columns=["Default_ind","States","tot_credit_debt",
    ↳"credit_age", "credit_good_age", "card_age","rep_income","inq_12_month_num"])
```

```

df_test = df_test.join(dummies)
df_test = df_test.join(default_ind)
util_vars = ["uti_card", "uti_50plus_pct", "uti_max_credit_line",
    ↪ "uti_card_50plus_pct"]
util_credit = df_test[util_vars]
util_credit["mean"] = util_credit.mean(axis=1)
df_test = df_test.drop(columns=["uti_50plus_pct", "uti_max_credit_line",
    ↪ "uti_card_50plus_pct"])
df_test["uti_card"] = util_credit["mean"]
df_test = df_test.rename(columns={"uti_card": "mean_uti"})
past_due_vars = ["non_mtg_acc_past_due_6_months_num",
    ↪ "mortgages_past_due_6_months_num"]
past_due_credit = df_test[past_due_vars]
past_due_credit["past_due_6_months_num"] =
    ↪ past_due_credit["non_mtg_acc_past_due_6_months_num"] +
    ↪ past_due_credit["mortgages_past_due_6_months_num"]
df_test = df_test.drop(columns=["mortgages_past_due_6_months_num"])
df_test["non_mtg_acc_past_due_6_months_num"] =
    ↪ past_due_credit["past_due_6_months_num"]
df_test = df_test.rename(columns={"non_mtg_acc_past_due_6_months_num":
    ↪ "past_due_6_months_num"})
df_test_norm = ((df_test - df_test.min()) / (df_test.max() - df_test.min())) * 20

```

<ipython-input-212-fd2b47dd7047>:8: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

util_credit["mean"] = util_credit.mean(axis=1)
<ipython-input-212-fd2b47dd7047>:14: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

past_due_credit["past_due_6_months_num"] =
past_due_credit["non_mtg_acc_past_due_6_months_num"] +
past_due_credit["mortgages_past_due_6_months_num"]

```

Check data:

```

[43]: print(df_test_norm.sample)
      print(df_test_norm.shape)
      print(df_test_norm.dtypes)

```

```

<bound method NDFrame.sample of          avg_card_debt
non_mtg_acc_past_due_12_months_num  \

```

0	0.723813	0.0
1	2.571408	0.0
2	2.007778	0.0
3	2.326841	0.0
4	2.071129	0.0
...	...	...
4995	1.534659	0.0
4996	2.550137	0.0
4997	1.506616	0.0
4998	1.193691	0.0
4999	2.871439	0.0

	past_due_6_months_num	credit_past_due_amount	card_inq_24_month_num	\
0	0.0	0.0	5.333333	
1	0.0	0.0	5.333333	
2	0.0	0.0	1.333333	
3	0.0	0.0	0.000000	
4	0.0	0.0	0.000000	
...	...	...	...	
4995	0.0	0.0	5.333333	
4996	0.0	0.0	6.666667	
4997	0.0	0.0	6.666667	
4998	0.0	0.0	2.666667	
4999	0.0	0.0	10.666667	

	card_open_36_month_num	auto_open_36_month_num	mean_util	ind_acc_XYZ	\
0	0.0	10.0	13.749481	0.0	
1	0.0	0.0	15.757954	20.0	
2	0.0	0.0	8.624625	0.0	
3	0.0	0.0	9.514111	20.0	
4	0.0	0.0	12.012403	0.0	
...	...	...	...	...	
4995	0.0	0.0	10.194068	0.0	
4996	0.0	0.0	6.574378	20.0	
4997	10.0	10.0	9.373143	20.0	
4998	0.0	0.0	0.907474	0.0	
4999	10.0	0.0	12.528196	0.0	

	AL	FL	GA	LA	MS	NC	SC	Default_ind
0	0.0	20.0	0.0	0.0	0.0	0.0	0.0	0.0
1	20.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	20.0	0.0	0.0	0.0
3	20.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	20.0	0.0	0.0	0.0	0.0	0.0
...	...	...	...	...	...	...	...	...
4995	0.0	0.0	0.0	20.0	0.0	0.0	0.0	0.0
4996	0.0	0.0	0.0	0.0	0.0	0.0	20.0	0.0
4997	0.0	0.0	0.0	20.0	0.0	0.0	0.0	0.0

```

4998    0.0    0.0    0.0    0.0  20.0    0.0    0.0          0.0
4999    0.0    0.0    0.0    0.0  20.0    0.0    0.0          0.0

```

```

[5000 rows x 17 columns]>
(5000, 17)
avg_card_debt                float64
non_mtg_acc_past_due_12_months_num  float64
past_due_6_months_num          float64
credit_past_due_amount         float64
card_inq_24_month_num          float64
card_open_36_month_num         float64
auto_open_36_month_num         float64
mean_util                      float64
ind_acc_XYZ                   float64
AL                             float64
FL                             float64
GA                             float64
LA                             float64
MS                             float64
NC                             float64
SC                             float64
Default_ind                   float64
dtype: object

```

Scoring test data:

```
[139]: lrModel.score(df_test_norm.iloc[:,range(0,16)], df_test_norm.iloc[:,16])
```

```
[139]: 0.9358
```

Testing with a confusion matrix:

```
[140]: confusion_matrix(df_test_norm.iloc[:,16],lrModel.predict(df_test_norm.iloc[:,
→,range(0,16)]))
```


```
[140]: array([[4552,  47],
              [ 274, 127]], dtype=int64)
```

Interpreting logistic regression model using ROC curve (2b):

```
[131]: lr_probs = lr.predict_proba(df_test_norm.iloc[:,range(0,16)])
lr_probs = lr_probs[:, 1]
lr_fpr, lr_tpr, _ = roc_curve(df_test_norm.iloc[:,16], lr_probs, pos_label = 20.
→0)
plt.plot(lr_fpr, lr_tpr, marker='.', label='Logistic Regression Classifier')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()

```





Find feature importance (2c):

A positive score indicates the greater the value of the feature, the more likely the person is to default, while a negative score indicates the greater the value of the feature, the less likely the person is to default.

A high absolute score indicates strong importance in model, while a low absolute score indicates weak importance in model.

Important Notes: - Having had an account in XYZ does not seem to impact the model by much - Average card debt, past due variables for 12 and 6 months, and utilization variables seem to have the greatest impact - Although the state variables seem to cause a great impact, in reality each observation must be one of them, and thus, the difference between them is the only thing that matters for interpretation. In this case, since the importance of all state variables is very close overall, it seems that states do not have a great deal of impact on model prediction

```
[179]: variables = []
for var in df_train_norm.columns:
    variables.append(var)
importance = lrModel.coef_[0]
index = 0
for i,v in enumerate(importance):
    print('Feature: {variable}\n Score: {v}\n\n'.format(variable = variables[index], v = v))
    index += 1
```

```
plt.bar([x for x in range(len(importance))], importance)
plt.show()
```

Feature: avg\_card\_debt  
Score: 0.1499596391270964

Feature: non\_mtg\_acc\_past\_due\_12\_months\_num  
Score: 0.23556759045910966

Feature: past\_due\_6\_months\_num  
Score: 0.1542133033538216

Feature: credit\_past\_due\_amount  
Score: -0.04906488819908006

Feature: card\_inq\_24\_month\_num  
Score: 0.08606432328090474

Feature: card\_open\_36\_month\_num  
Score: 0.016807004273557042

Feature: auto\_open\_36\_month\_num  
Score: 0.006285595831829881

Feature: mean\_util  
Score: 0.27185106310181156

Feature: ind\_acc\_XYZ  
Score: -0.013710761590003758

Feature: AL  
Score: -0.3130380831427379

Feature: FL  
Score: -0.317889156327858

Feature: GA

Score: -0.318481497908399

Feature: LA

Score: -0.31415871337745066

Feature: MS


Score: -0.3113582401283215

Feature: NC

Score: -0.31565900138163133

Feature: SC

Score: -0.3103977152431603



## 5 Data Preprocessing - Random Forests

We will reimport the data since we will be using a different approach to data preprocessing in preparation of random forest model.

```
[144]: df_test = pd.read_csv('Simulated_Data_Test.csv')
df_train = pd.read_csv('Simulated_Data_Train.csv')
df_validation = pd.read_csv('Simulated_Data_Validation.csv')
```

Based on EDA, we have determined which explanatory variables need to be dropped: rep\_income, credit\_age, credit\_good\_age, and card\_age.

```
[145]: df_train = df_train.drop(columns=["credit_age", "credit_good_age",
↳ "card_age", "rep_income"])
```

Next, let's create dummy variables for states

```
[146]: dummies = pd.get_dummies(df_train['States'])
default_ind = df_train['Default_ind']
df_train = df_train.drop(columns=['States', 'Default_ind'])
df_train = df_train.join(dummies)
df_train = df_train.join(default_ind)
```

We need to get rid of NA values next. let's check how many there are:

```
[147]: print(df_train.isnull().sum())
```

```
tot_credit_debt          0
avg_card_debt            0
non_mtg_acc_past_due_12_months_num    0
non_mtg_acc_past_due_6_months_num    0
mortgages_past_due_6_months_num    0
credit_past_due_amount    0
inq_12_month_num        0
card_inq_24_month_num    0
card_open_36_month_num    0
auto_open_36_month_num    0
uti_card                0
uti_50plus_pct          0
uti_max_credit_line      0
uti_card_50plus_pct      2055
ind_acc_XYZ             0
AL                      0
FL                      0
GA                      0
LA                      0
MS                      0
NC                      0
SC                      0
Default_ind             0
dtype: int64
```

We want to preserve the utilization features as we are not concerned with collinearity for this model. Thus, we have two options: - Remove observations or column with missing values. - Impute missing values for uti\_card\_50plus\_pct

Removing observations will rid us of 10% of our data, which is not ideal. Removing `uti_card_50plus_pct` is detrimental since we will be getting rid of a feature which is highly correlated with the response variable. Thus we decide to impute the missing values of `uti_card_50plus_pct`.

We know `uti_card_50plus_pct` is highly collinear with the other utilization features. Therefore, we will be using simple mean imputation with the utilization variables.

```
[148]: imputer = SingleImputer(  
        strategy='mean',  
        predictors=['uti_max_credit_line', 'uti_50plus_pct', 'uti_card'],  
        seed = 100)  
        imputed_df_train = imputer.fit_transform(df_train)
```

Check if imputation was succesful:

```
[149]: print(imputed_df_train.isnull().sum())
```

```
tot_credit_debt          0  
avg_card_debt            0  
non_mtg_acc_past_due_12_months_num  0  
non_mtg_acc_past_due_6_months_num  0  
mortgages_past_due_6_months_num  0  
credit_past_due_amount    0  
inq_12_month_num         0  
card_inq_24_month_num     0  
card_open_36_month_num    0  
auto_open_36_month_num    0  
uti_card                 0  
uti_50plus_pct           0  
uti_max_credit_line       0  
uti_card_50plus_pct       0  
ind_acc_XYZ              0  
AL                       0  
FL                       0  
GA                       0  
LA                       0  
MS                       0  
NC                       0  
SC                       0  
Default_ind              0  
dtype: int64
```

Check data:

Everything looks ready for the model.

```
[150]: print(imputed_df_train.sample)  
        print(imputed_df_train.shape)
```

```
print(imputed_df_train.dtypes)
```

```
<bound method NDFrame.sample of
non_mtg_acc_past_due_12_months_num \
0          80826.71          15872.99          0.0
1          96052.60          12178.02          0.0
2          75212.76          12052.24          0.0
3          70727.84           8416.80          2.0
4          41604.47          10611.97          0.0
...
19995        104765.01          13905.40          0.0
19996        83990.07          10325.02          0.0
19997        107606.69          17838.79          1.0
19998        78787.72          11447.61          0.0
19999        78296.90          10053.16          0.0

non_mtg_acc_past_due_6_months_num  mortgages_past_due_6_months_num \
0                                0.0                                0.0
1                                0.0                                0.0
2                                0.0                                0.0
3                                1.0                                1.0
4                                0.0                                0.0
...
19995                            0.0                                0.0
19996                            0.0                                0.0
19997                            0.0                                0.0
19998                            0.0                                0.0
19999                            0.0                                0.0

credit_past_due_amount  inq_12_month_num  card_inq_24_month_num \
0                    0.00                3.0                4.0
1                    0.00                2.0                4.0
2                    0.00                1.0                3.0
3                11013.96                0.0                1.0
4                    0.00                0.0                2.0
...
19995                    0.00                4.0                4.0
19996                    0.00                0.0                1.0
19997                    0.00                3.0                3.0
19998                    0.00                1.0                2.0
19999                    0.00                2.0                4.0

card_open_36_month_num  auto_open_ 36_month_num  ... \
0                    0.0                    0.0  ...
1                    1.0                    0.0  ...
2                    0.0                    1.0  ...
3                    1.0                    0.0  ...
4                    0.0                    0.0  ...
```

```

...
19995      0.0      0.0 ...
19996      0.0      0.0 ...
19997      0.0      1.0 ...
19998      0.0      0.0 ...
19999      0.0      0.0 ...

```

```

      uti_card_50plus_pct  ind_acc_XYZ  AL  FL  GA  LA  MS  NC  SC  \
0      0.489594      0.0    1    0    0    0    0    0    0
1      0.587351      0.0    0    1    0    0    0    0    0
2      0.413293      0.0    1    0    0    0    0    0    0
3      0.466810      0.0    0    0    0    0    0    0    1
4      0.588442      1.0    0    0    0    1    0    0    0
...
19995      0.481113      1.0    0    0    1    0    0    0    0
19996      0.489594      0.0    0    0    0    0    0    1    0
19997      0.463120      0.0    0    0    0    0    0    0    1
19998      0.281647      0.0    1    0    0    0    0    0    0
19999      0.512513      0.0    1    0    0    0    0    0    0

```

```

      Default_ind
0      0.0
1      0.0
2      0.0
3      1.0
4      0.0
...
19995      0.0
19996      0.0
19997      0.0
19998      0.0
19999      0.0

```

[20000 rows x 23 columns]>

(20000, 23)

```

tot_credit_debt      float64
avg_card_debt      float64
non_mtg_acc_past_due_12_months_num      float64
non_mtg_acc_past_due_6_months_num      float64
mortgages_past_due_6_months_num      float64
credit_past_due_amount      float64
inq_12_month_num      float64
card_inq_24_month_num      float64
card_open_36_month_num      float64
auto_open_36_month_num      float64
uti_card      float64
uti_50plus_pct      float64
uti_max_credit_line      float64

```

```

uti_card_50plus_pct          float64
ind_acc_XYZ                  float64
AL                            uint8
FL                            uint8
GA                            uint8
LA                            uint8
MS                            uint8
NC                            uint8
SC                            uint8
Default_ind                  float64
dtype: object

```

## 6 Model Building and Testing - Random Forests

We will firstly preprocess the validation data set in same way as train set to prepare for hyperparameter tuning.

```

[151]: df_validation = df_validation.drop(columns=["credit_age", "credit_good_age", "
        ↪ "card_age", "rep_income"])
        dummies = pd.get_dummies(df_validation['States'])
        default_ind = df_validation['Default_ind']
        df_validation = df_validation.drop(columns=['States', 'Default_ind'])
        df_validation = df_validation.join(dummies)
        df_validation = df_validation.join(default_ind)
        imputed_df_validation = imputer.fit_transform(df_validation)

```

Check validation data:

Everything looks in order.

```

[152]: print(imputed_df_validation.sample)
        print(imputed_df_validation.shape)
        print(imputed_df_validation.dtypes)

```

```

<bound method NDFrame.sample of          tot_credit_debt  avg_card_debt
non_mtg_acc_past_due_12_months_num  \
0          63651.27          9019.99          0.0
1         105559.29         16692.19          0.0
2          96062.99         10509.13          0.0
3          84417.40         13873.96          0.0
4         100623.91         15592.09          0.0
...          ...          ...          ...
2995         90748.88         11481.81          0.0
2996         101930.98         15242.30          0.0
2997          74738.73         12175.49          0.0
2998         120357.58         14477.70          0.0
2999          74240.17         11873.45          0.0

```



	non_mtg_acc_past_due_6_months_num	mortgages_past_due_6_months_num	\
0	0.0	0.0	
1	0.0	0.0	
2	0.0	0.0	
3	0.0	0.0	
4	0.0	0.0	
...	...	...	
2995	0.0	0.0	
2996	0.0	0.0	
2997	0.0	0.0	
2998	0.0	0.0	
2999	0.0	0.0	

	credit_past_due_amount	inq_12_month_num	card_inq_24_month_num	\
0	0.0	0.0	0.0	
1	0.0	4.0	7.0	
2	0.0	2.0	3.0	
3	0.0	1.0	1.0	
4	0.0	6.0	9.0	
...	...	...	...	
2995	0.0	0.0	0.0	
2996	0.0	2.0	0.0	
2997	0.0	2.0	4.0	
2998	0.0	2.0	2.0	
2999	0.0	4.0	7.0	

	card_open_36_month_num	auto_open_36_month_num	...	\
0	0.0	0.0	...	
1	0.0	1.0	...	
2	0.0	0.0	...	
3	0.0	0.0	...	
4	0.0	0.0	...	
...	...	...	...	
2995	0.0	1.0	...	
2996	0.0	1.0	...	
2997	0.0	0.0	...	
2998	0.0	1.0	...	
2999	1.0	1.0	...	

	uti_card_50plus_pct	ind_acc_XYZ	AL	FL	GA	LA	MS	NC	SC	\
0	0.530109	0.0	0	1	0	0	0	0	0	
1	0.388792	0.0	0	0	0	0	1	0	0	
2	0.450523	0.0	0	0	0	0	1	0	0	
3	0.562153	0.0	0	0	0	0	1	0	0	
4	0.576409	1.0	0	0	0	0	0	1	0	
...	...	...	..	..	..	..	..	..	..	
2995	0.559056	0.0	0	1	0	0	0	0	0	
2996	0.564803	0.0	1	0	0	0	0	0	0	

2997	0.651801	0.0	0	0	0	1	0	0	0
2998	0.557991	1.0	0	0	0	0	0	1	0
2999	0.736850	0.0	0	0	0	0	1	0	0

	Default_ind
0	1.0
1	0.0
2	0.0
3	0.0
4	0.0
...	...
2995	0.0
2996	0.0
2997	0.0
2998	0.0
2999	0.0

```
[3000 rows x 23 columns]>
(3000, 23)
tot_credit_debt          float64
avg_card_debt            float64
non_mtg_acc_past_due_12_months_num float64
non_mtg_acc_past_due_6_months_num  float64
mortgages_past_due_6_months_num    float64
credit_past_due_amount    float64
inq_12_month_num          float64
card_inq_24_month_num     float64
card_open_36_month_num    float64
auto_open_36_month_num    float64
uti_card                  float64
uti_50plus_pct            float64
uti_max_credit_line       float64
uti_card_50plus_pct       float64
ind_acc_XYZ               float64
AL                         uint8
FL                         uint8
GA                         uint8
LA                         uint8
MS                         uint8
NC                         uint8
SC                         uint8
Default_ind               float64
dtype: object
```

We will now perform hyperparamter tuning using our validation data set (2a).

The parameters we have chosen are: - Number of estimators: 10, 100, 1000 - Max features: square root, log base 2

```
[90]: model = RandomForestClassifier()
n_estimators = [10, 100, 1000]
max_features = ['sqrt', 'log2']
grid = dict(n_estimators=n_estimators,max_features=max_features)
cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
grid_search = GridSearchCV(estimator=model, param_grid=grid, n_jobs=-1, cv=cv,
    ↳scoring='accuracy',error_score=0)
grid_result = grid_search.fit(imputed_df_validation.iloc[:,range(0,22)],
    ↳imputed_df_validation.iloc[:,22])
print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))
```

```
Best: 0.941111 using {'max_features': 'log2', 'n_estimators': 100}
0.936222 (0.009574) with: {'max_features': 'sqrt', 'n_estimators': 10}
0.940667 (0.010306) with: {'max_features': 'sqrt', 'n_estimators': 100}
0.940556 (0.009929) with: {'max_features': 'sqrt', 'n_estimators': 1000}
0.937778 (0.009480) with: {'max_features': 'log2', 'n_estimators': 10}
0.941111 (0.010830) with: {'max_features': 'log2', 'n_estimators': 100}
0.940444 (0.010532) with: {'max_features': 'log2', 'n_estimators': 1000}
```

Based on this, the most accurate model has the following hyperparameters: - Number of estimators: 100 - Max features: log base 2

Train model based on chosen hyperparameters using the train data set:

```
[153]: rfModel = RandomForestClassifier(max_features="log2",n_estimators=100)
rfModel = rfModel.fit(imputed_df_train.iloc[:,range(0,22)], imputed_df_train.
    ↳iloc[:,22])
```

Transforming test data to match training data:

```
[154]: df_test = df_test.drop(columns=["credit_age", "credit_good_age",
    ↳"card_age", "rep_income"])
dummies = pd.get_dummies(df_test['States'])
default_ind = df_test['Default_ind']
df_test = df_test.drop(columns=['States', 'Default_ind'])
df_test = df_test.join(dummies)
df_test = df_test.join(default_ind)
imputed_df_test = imputer.fit_transform(df_test)
```

Check data:

```
[155]: print(imputed_df_test.sample)
print(imputed_df_test.shape)
print(imputed_df_test.dtypes)
```

```

<bound method NDFrame.sample of
non_mtg_acc_past_due_12_months_num \
0          40477.81          7766.64          0.0
1          106760.98          16606.98          0.0
2          121428.34          13910.13          0.0
3           96515.05          15436.78          0.0
4          123760.22          14213.25          0.0
...
4995         90788.44          11646.36          0.0
4996         91052.81          16505.20          0.0
4997         71061.71          11512.18          0.0
4998         82162.45          10014.90          0.0
4999        116943.32          18042.56          0.0

```

```

non_mtg_acc_past_due_6_months_num mortgages_past_due_6_months_num \
0          0.0          0.0
1          0.0          0.0
2          0.0          0.0
3          0.0          0.0
4          0.0          0.0
...
4995         0.0          0.0
4996         0.0          0.0
4997         0.0          0.0
4998         0.0          0.0
4999         0.0          0.0

```

```

credit_past_due_amount inq_12_month_num card_inq_24_month_num \
0          0.0          2.0          4.0
1          0.0          2.0          4.0
2          0.0          0.0          1.0
3          0.0          0.0          0.0
4          0.0          0.0          0.0
...
4995         0.0          0.0          4.0
4996         0.0          3.0          5.0
4997         0.0          3.0          5.0
4998         0.0          1.0          2.0
4999         0.0          4.0          8.0

```

```

card_open_36_month_num auto_open_36_month_num ... \
0          0.0          1.0 ...
1          0.0          0.0 ...
2          0.0          0.0 ...
3          0.0          0.0 ...
4          0.0          0.0 ...
...
4995         0.0          0.0 ...

```

4996	0.0	0.0	...
4997	1.0	1.0	...
4998	0.0	0.0	...
4999	1.0	0.0	...

	uti_card_50plus_pct	ind_acc_XYZ	AL	FL	GA	LA	MS	NC	SC	\
0	0.633699	0.0	0	1	0	0	0	0	0	
1	0.746901	1.0	1	0	0	0	0	0	0	
2	0.458438	0.0	0	0	0	0	1	0	0	
3	0.582352	1.0	1	0	0	0	0	0	0	
4	0.577697	0.0	0	0	1	0	0	0	0	
...	...	...	..	..	..	..	..	..	..	
4995	0.398411	0.0	0	0	0	1	0	0	0	
4996	0.286623	1.0	0	0	0	0	0	0	1	
4997	0.443688	1.0	0	0	0	1	0	0	0	
4998	0.178131	0.0	0	0	0	0	1	0	0	
4999	0.503552	0.0	0	0	0	0	1	0	0	

	Default_ind
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0
...	...
4995	0.0
4996	0.0
4997	0.0
4998	0.0
4999	0.0

[5000 rows x 23 columns]>

(5000, 23)

tot_credit_debt	float64
avg_card_debt	float64
non_mtg_acc_past_due_12_months_num	float64
non_mtg_acc_past_due_6_months_num	float64
mortgages_past_due_6_months_num	float64
credit_past_due_amount	float64
inq_12_month_num	float64
card_inq_24_month_num	float64
card_open_36_month_num	float64
auto_open_36_month_num	float64
uti_card	float64
uti_50plus_pct	float64
uti_max_credit_line	float64
uti_card_50plus_pct	float64
ind_acc_XYZ	float64

AL	uint8
FL	uint8
GA	uint8
LA	uint8
MS	uint8
NC	uint8
SC	uint8
Default_ind	float64
dtype:	object

Scoring test data:

```
[156]: rfModel.score(imputed_df_test.iloc[:,range(0,22)], imputed_df_test.iloc[:,22])
```

```
[156]: 0.9366
```


Testing with a confusion matrix:

```
[157]: confusion_matrix(imputed_df_test.iloc[:,22],rfModel.predict(imputed_df_test.
↪iloc[:,range(0,22)]))
```

```
[157]: array([[4568,  31],
              [ 286, 115]], dtype=int64)
```

ROC curve for random forest classifier (2b):

```
[158]: rf_probs = rfModel.predict_proba(imputed_df_test.iloc[:,range(0,22)])
rf_probs = rf_probs[:, 1]
rf_fpr, rf_tpr, _ = roc_curve(imputed_df_test.iloc[:,22], rf_probs)
plt.plot(rf_fpr, rf_tpr, marker='.', label='Random Forest Classifier')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```



Find feature importance (2c):

The higher the score, the more important the feature.

Important Notes: - Despite not presenting much correlation with response variable, total credit debt and average card debt play the most important roles in determining defaults - Utilization variables are all very important - State, XYZ previous account indicator, and credit products opened in 36 months variables play a marginal role

```
[182]: variables = []
for var in imputed_df_train.columns:
    variables.append(var)
importance = rfModel.feature_importances_
index = 0
for i,v in enumerate(importance):
    print('Feature: {variable}\n Score: {v}\n\n'.format(variable = variables[index], v = v))
    index += 1
plt.bar([x for x in range(len(importance))], importance)
plt.show()
```

```
Feature: tot_credit_debt
Score: 0.10116110103522905
```

Feature: avg\_card\_debt  
Score: 0.15612549704029433

Feature: non\_mtg\_acc\_past\_due\_12\_months\_num  
Score: 0.051588433741830204

Feature: non\_mtg\_acc\_past\_due\_6\_months\_num  
Score: 0.026114066030776564

Feature: mortgages\_past\_due\_6\_months\_num  
Score: 0.04471094287968974

Feature: credit\_past\_due\_amount  
Score: 0.07272974478355873

Feature: inq\_12\_month\_num  
Score: 0.040155479644069965

Feature: card\_inq\_24\_month\_num  
Score: 0.0476335411984235

Feature: card\_open\_36\_month\_num  
Score: 0.012521132722843388

Feature: auto\_open\_36\_month\_num  
Score: 0.009328620784239058

Feature: uti\_card  
Score: 0.10270683546384929

Feature: uti\_50plus\_pct  
Score: 0.08696333953190324

Feature: uti\_max\_credit\_line  
Score: 0.09047046122791504



Feature: uti\_card\_50plus\_pct  
Score: 0.08913573380667056

Feature: ind\_acc\_XYZ  
Score: 0.011392454240505133

Feature: AL  
Score: 0.008526349581013072

Feature: FL  
Score: 0.00890519489911215


Feature: GA  
Score: 0.007348783410222652

Feature: LA  
Score: 0.008199211337674586

Feature: MS  
Score: 0.008037655223021628

Feature: NC  
Score: 0.008202140317150004

Feature: SC  
Score: 0.008043281100008147




## 7 Comparing Models

Comparing ROC curves and AUC (2b):

Random forest performs better, as expected, because the true positive rate was better. This is confirmed by the AUC calculation as well.

```
[229]: fig, ax = plt.subplots(figsize=(12,10))
plt.plot(rf_fpr, rf_tpr, marker='.', label='Random Forest Classifier')
plt.plot(lr_fpr, lr_tpr, marker='.', label='Logistic Regression Classifier')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```



```
[136]: lr_auc = roc_auc_score(df_test_norm.iloc[:,16], lr_probs)
       rf_auc = roc_auc_score(imputed_df_test.iloc[:,22], rf_probs)
       print('Logistic Regression: ROC AUC=%.3f' % (lr_auc))
       print('Random Forest: ROC AUC=%.3f' % (rf_auc))
```

Logistic Regression: ROC AUC=0.806

Random Forest: ROC AUC=0.852

Comparing Learning curves (3): - While the training score for the random forest model is very high (a natural consequence of an algorithm that is extremely robust), the cross validation scores are very comparable throughout - Unsurprisingly, logistic regression is found to be much more scalable - Model performance is slightly higher for random forest model

```
[161]: def plot_learning_curve(estimator, title, X, y, axes=None, ylim=None, cv=None,
                             n_jobs=None, train_sizes=np.linspace(.1, 1.0, 5)):
       if axes is None:
           _, axes = plt.subplots(1, 3, figsize=(20, 5))

       axes[0].set_title(title)
```

```

if ylim is not None:
    axes[0].set_ylim(*ylim)
axes[0].set_xlabel("Training examples")
axes[0].set_ylabel("Score")

train_sizes, train_scores, test_scores, fit_times, _ = \
    learning_curve(estimator, X, y, cv=cv, n_jobs=n_jobs,
                    train_sizes=train_sizes,
                    return_times=True)
train_scores_mean = np.mean(train_scores, axis=1)
train_scores_std = np.std(train_scores, axis=1)
test_scores_mean = np.mean(test_scores, axis=1)
test_scores_std = np.std(test_scores, axis=1)
fit_times_mean = np.mean(fit_times, axis=1)
fit_times_std = np.std(fit_times, axis=1)

axes[0].grid()
axes[0].fill_between(train_sizes, train_scores_mean - train_scores_std,
                     train_scores_mean + train_scores_std, alpha=0.1,
                     color="r")
axes[0].fill_between(train_sizes, test_scores_mean - test_scores_std,
                     test_scores_mean + test_scores_std, alpha=0.1,
                     color="g")
axes[0].plot(train_sizes, train_scores_mean, 'o-', color="r",
              label="Training score")
axes[0].plot(train_sizes, test_scores_mean, 'o-', color="g",
              label="Cross-validation score")
axes[0].legend(loc="best")

axes[1].grid()
axes[1].plot(train_sizes, fit_times_mean, 'o-')
axes[1].fill_between(train_sizes, fit_times_mean - fit_times_std,
                     fit_times_mean + fit_times_std, alpha=0.1)
axes[1].set_xlabel("Training examples")
axes[1].set_ylabel("fit_times")
axes[1].set_title("Scalability of the model")

axes[2].grid()
axes[2].plot(fit_times_mean, test_scores_mean, 'o-')
axes[2].fill_between(fit_times_mean, test_scores_mean - test_scores_std,
                     test_scores_mean + test_scores_std, alpha=0.1)
axes[2].set_xlabel("fit_times")
axes[2].set_ylabel("Score")
axes[2].set_title("Performance of the model")

return plt

```

```

fig, axes = plt.subplots(3, 2, figsize=(10, 15))







X, y = df_train_norm.iloc[:,range(0,16)], df_train_norm.iloc[:,16]
title = "Learning Curves (Logistic Regression)"
cv = ShuffleSplit(n_splits=100, test_size=0.2, random_state=0)

plot_learning_curve(lrModel, title, X, y, axes=axes[:, 0], ylim=(0.7, 1.01),
                    cv=cv, n_jobs=4)

X, y = imputed_df_train.iloc[:,range(0,22)], imputed_df_train.iloc[:,22]
title = "Learning Curves (Random Forest)"
cv = ShuffleSplit(n_splits=50, test_size=0.2, random_state=0)
plot_learning_curve(rfModel, title, X, y, axes=axes[:, 1], ylim=(0.7, 1.01),
                    cv=cv, n_jobs=4)

plt.show()

```



## 8 Code References

*Thank you to the following:*

(1) Aniruddha Bhandari:

[What is Multicollinearity? Here's Everything You Need to Know](#)

(2) Jason Brownlee, PhD of Machine Learning Mastery:

a. [Tune Hyperparameters for Classification Machine Learning Algorithms](#)

b. [How to Use ROC Curves and Precision-Recall Curves for Classification in Python](#)

c. [How to Calculate Feature Importance With Python](#)

(3) SKLearn:

[Plotting Learning Curves](#)